

ASSIGNMENT 02

NAME: BHOSALE SRUSHTI AMAR

CLASS: TY AIDS-A

ROLL NO.: 23107008

```
In [5]: import pandas as pd  
import numpy as np
```

```
In [3]: df=pd.read_csv("Downloads/WA_Fn-UseC_-HR-Employee-Attrition.csv")
```

```
In [4]: df
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Edu
0	41	Yes	Travel_Rarely	1102	Sales		1
1	49	No	Travel_Frequently	279	Research & Development		8
2	37	Yes	Travel_Rarely	1373	Research & Development		2
3	33	No	Travel_Frequently	1392	Research & Development		3
4	27	No	Travel_Rarely	591	Research & Development		2
...
1465	36	No	Travel_Frequently	884	Research & Development		23
1466	39	No	Travel_Rarely	613	Research & Development		6
1467	27	No	Travel_Rarely	155	Research & Development		4
1468	49	No	Travel_Frequently	1023	Sales		2
1469	34	No	Travel_Rarely	628	Research & Development		8

1470 rows × 35 columns

```
In [6]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
 #   Column           Non-Null Count Dtype  
--- 
 0   Age              1470 non-null   int64   
 1   Attrition        1470 non-null   object  
 2   BusinessTravel   1470 non-null   object  
 3   DailyRate         1470 non-null   int64   
 4   Department        1470 non-null   object  
 5   DistanceFromHome 1470 non-null   int64   
 6   Education         1470 non-null   int64   
 7   EducationField    1470 non-null   object  
 8   EmployeeCount     1470 non-null   int64   
 9   EmployeeNumber    1470 non-null   int64   
 10  EnvironmentSatisfaction 1470 non-null   int64   
 11  Gender            1470 non-null   object  
 12  HourlyRate        1470 non-null   int64   
 13  JobInvolvement   1470 non-null   int64   
 14  JobLevel          1470 non-null   int64   
 15  JobRole           1470 non-null   object  
 16  JobSatisfaction   1470 non-null   int64   
 17  MaritalStatus     1470 non-null   object  
 18  MonthlyIncome     1470 non-null   int64   
 19  MonthlyRate       1470 non-null   int64   
 20  NumCompaniesWorked 1470 non-null   int64   
 21  Over18            1470 non-null   object  
 22  Overtime          1470 non-null   object  
 23  PercentSalaryHike 1470 non-null   int64   
 24  PerformanceRating 1470 non-null   int64   
 25  RelationshipSatisfaction 1470 non-null   int64   
 26  StandardHours     1470 non-null   int64   
 27  StockOptionLevel   1470 non-null   int64   
 28  TotalWorkingYears 1470 non-null   int64   
 29  TrainingTimesLastYear 1470 non-null   int64   
 30  WorkLifeBalance   1470 non-null   int64   
 31  YearsAtCompany    1470 non-null   int64   
 32  YearsInCurrentRole 1470 non-null   int64   
 33  YearsSinceLastPromotion 1470 non-null   int64   
 34  YearsWithCurrManager 1470 non-null   int64 

dtypes: int64(26), object(9)
memory usage: 402.1+ KB

```

In [10]: `df.tail()`

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Edu
1465	36	No	Travel_Frequently	884	Research & Development		23
1466	39	No	Travel_Rarely	613	Research & Development		6
1467	27	No	Travel_Rarely	155	Research & Development		4
1468	49	No	Travel_Frequently	1023	Sales		2
1469	34	No	Travel_Rarely	628	Research & Development		8

5 rows × 35 columns

In [11]: `df.isnull().sum()`

```
Out[11]: Age          0
Attrition      0
BusinessTravel 0
DailyRate       0
Department      0
DistanceFromHome 0
Education        0
EducationField   0
EmployeeCount    0
EmployeeNumber   0
EnvironmentSatisfaction 0
Gender          0
HourlyRate      0
JobInvolvement 0
JobLevel         0
JobRole          0
JobSatisfaction 0
MaritalStatus    0
MonthlyIncome     0
MonthlyRate      0
NumCompaniesWorked 0
Over18           0
OverTime          0
PercentSalaryHike 0
PerformanceRating 0
RelationshipSatisfaction 0
StandardHours    0
StockOptionLevel 0
TotalWorkingYears 0
TrainingTimesLastYear 0
WorkLifeBalance 0
YearsAtCompany    0
YearsInCurrentRole 0
YearsSinceLastPromotion 0
YearsWithCurrManager 0
dtype: int64
```

```
In [12]: df = df.dropna()
```

```
In [19]: print(df.columns.tolist())
```

```
['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department', 'DistanceFromHome', 'Education', 'EducationField', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'Job Involvement', 'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'OverTime', 'PercentSalaryHike', 'Performance Rating', 'RelationshipSatisfaction', 'StockOptionLevel', 'TotalWorkingYears', 'Trainin gTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinc eLastPromotion', 'YearsWithCurrManager']
```

```
In [20]: num_cols = df.select_dtypes(include=['int64', 'float64']).columns
```

```
In [22]: Q1 = df[num_cols].quantile(0.25)
Q3 = df[num_cols].quantile(0.75)
IQR = Q3 - Q1
df1 = df[~((df[num_cols] < (Q1 - 1.5 * IQR)) |
            (df[num_cols] > (Q3 + 1.5 * IQR))).any(axis=1)]
```

```
In [23]: print("Before outlier removal:", df.shape)
print("After outlier removal (df1):", df1.shape)
```

```
Before outlier removal: (1470, 31)
After outlier removal (df1): (779, 31)
```

```
In [24]: from sklearn.preprocessing import LabelEncoder
```

```
In [25]: le = LabelEncoder()
```

```
In [26]: for col in df1.select_dtypes(include='object').columns:  
    df1[col] = le.fit_transform(df1[col])
```

/tmp/ipykernel_5149/2731492867.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    df1[col] = le.fit_transform(df1[col])
```

/tmp/ipykernel_5149/2731492867.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    df1[col] = le.fit_transform(df1[col])
```

/tmp/ipykernel_5149/2731492867.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    df1[col] = le.fit_transform(df1[col])
```

/tmp/ipykernel_5149/2731492867.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    df1[col] = le.fit_transform(df1[col])
```

/tmp/ipykernel_5149/2731492867.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    df1[col] = le.fit_transform(df1[col])
```

/tmp/ipykernel_5149/2731492867.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    df1[col] = le.fit_transform(df1[col])
```

/tmp/ipykernel_5149/2731492867.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    df1[col] = le.fit_transform(df1[col])
```

/tmp/ipykernel_5149/2731492867.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    df1[col] = le.fit_transform(df1[col])
```

```
In [27]: X = df1.drop('Attrition', axis=1)  
y = df1['Attrition']
```

```
In [28]: from sklearn.model_selection import train_test_split
```

```
In [31]: X_train, X_test, y_train, y_test = train_test_split(  
    X, y,  
    test_size=0.2,  
    random_state=42)
```

```
In [32]: from sklearn.preprocessing import StandardScaler
```

```
In [33]: scaler = StandardScaler()
```

```
In [34]: X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

```
In [35]: from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense
```

```
2026-02-05 10:27:35.031758: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:  
485] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT  
when one has already been registered  
2026-02-05 10:27:35.218667: E external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:  
8454] Unable to register cuDNN factory: Attempting to register factory for plugin cuDN  
N when one has already been registered  
2026-02-05 10:27:35.339587: E external/local_xla/xla/stream_executor/cuda/cuda_blas.c  
c:1452] Unable to register cuBLAS factory: Attempting to register factory for plugin c  
uBLAS when one has already been registered  
2026-02-05 10:27:35.928529: I tensorflow/core/platform/cpu_feature_guard.cc:210] This  
TensorFlow binary is optimized to use available CPU instructions in performance-critic  
al operations.  
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlo  
w with the appropriate compiler flags.  
2026-02-05 10:27:37.883446: W tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF  
-TRT Warning: Could not find TensorRT
```

```
In [36]: model = Sequential()
```

```
In [37]: model.add(Dense(32, activation='relu', input_shape=(X_train.shape[1],)))  
model.add(Dense(16, activation='relu'))  
model.add(Dense(1, activation='sigmoid'))
```

```
/home/admin1/anaconda3/lib/python3.12/site-packages/keras/src/layers/core/dense.py:87:  
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using  
Sequential models, prefer using an `Input(shape)` object as the first layer in the mod  
el instead.  
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [38]: model.compile(  
    optimizer='adam',  
    loss='binary_crossentropy',  
    metrics=['accuracy'])
```

```
In [39]: model.fit(  
    X_train,  
    y_train,  
    epochs=20,  
    batch_size=16,  
    verbose=1)
```

```
Epoch 1/20
39/39 1s 1ms/step - accuracy: 0.6369 - loss: 0.6463
Epoch 2/20
39/39 0s 1ms/step - accuracy: 0.8145 - loss: 0.5069
Epoch 3/20
39/39 0s 1ms/step - accuracy: 0.8345 - loss: 0.4293
Epoch 4/20
39/39 0s 1ms/step - accuracy: 0.8338 - loss: 0.4041
Epoch 5/20
39/39 0s 1ms/step - accuracy: 0.8576 - loss: 0.3587
Epoch 6/20
39/39 0s 973us/step - accuracy: 0.8523 - loss: 0.3550
Epoch 7/20
39/39 0s 1ms/step - accuracy: 0.8692 - loss: 0.3342
Epoch 8/20
39/39 0s 1ms/step - accuracy: 0.8553 - loss: 0.3717
Epoch 9/20
39/39 0s 1ms/step - accuracy: 0.8770 - loss: 0.3027
Epoch 10/20
39/39 0s 1ms/step - accuracy: 0.8668 - loss: 0.3282
Epoch 11/20
39/39 0s 1ms/step - accuracy: 0.8957 - loss: 0.2755
Epoch 12/20
39/39 0s 1ms/step - accuracy: 0.8956 - loss: 0.2751
Epoch 13/20
39/39 0s 2ms/step - accuracy: 0.9047 - loss: 0.2580
Epoch 14/20
39/39 0s 1ms/step - accuracy: 0.8938 - loss: 0.2737
Epoch 15/20
39/39 0s 2ms/step - accuracy: 0.8902 - loss: 0.2829
Epoch 16/20
39/39 0s 2ms/step - accuracy: 0.8942 - loss: 0.2633
Epoch 17/20
39/39 0s 1ms/step - accuracy: 0.9035 - loss: 0.2554
Epoch 18/20
39/39 0s 1ms/step - accuracy: 0.9065 - loss: 0.2313
Epoch 19/20
39/39 0s 2ms/step - accuracy: 0.8824 - loss: 0.2524
Epoch 20/20
39/39 0s 1ms/step - accuracy: 0.9070 - loss: 0.2305
```

```
Out[39]: <keras.src.callbacks.history.History at 0x7f35b4105ac0>
```

```
In [43]: from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
y_pred = (model.predict(X_test) > 0.5).astype(int)
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
print("Confusion Matrix:", confusion_matrix(y_test, y_pred))
```

```
5/5 0s 2ms/step
Accuracy: 0.8461538461538461
      precision    recall  f1-score   support
          0       0.88     0.95     0.91      131
          1       0.53     0.32     0.40       25
   accuracy                           0.85      156
  macro avg       0.71     0.63     0.66      156
weighted avg       0.82     0.85     0.83      156
```

```
Confusion Matrix: [[124  7]
 [ 17  81]]
```

```
In [ ]:
```